

Ratnam_Project_2.R

dubey

Thu Apr 06 15:46:25 2017

#For each of the following series (from the fma package), make a graph of the data.

#If transforming seems appropriate, do so and describe the effect.

#Question 1) Monthly total of people on unemployed benefits in Australia (January 1956-July 1992).

#Loading the library

library(fma)

Warning: package 'fma' was built under R version 3.3.3

Loading required package: forecast

Warning: package 'forecast' was built under R version 3.3.3

#Loading the Package

dole

##	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
## 1956	4742	6128	6494	5379	6011	7003	9164	10333	9614	9545
## 1957	15711	13135	13077	15453	15995	18071	20291	20175	18975	17928
## 1958	29856	26879	24485	27745	27282	29418	29908	29278	26002	23826
## 1959	31486	28207	27669	27559	27924	27528	27410	24887	21904	19598
## 1960	23781	20020	18177	17732	16765	16310	14897	12940	11465	10364
## 1961	19257	20941	29718	35025	45110	57154	61499	62090	59561	48531
## 1962	56755	49740	45870	49136	47256	46324	45453	42333	36851	33952
## 1963	46178	40482	36394	37142	36424	38188	37174	31869	26575	21758
## 1964	28649	24226	21955	19937	18287	18129	17072	14924	12491	11160
## 1965	15831	13698	12111	12690	12585	12855	12137	10977	9993	9614
## 1966	19490	17611	16206	17560	18082	19482	19200	18918	17375	16122
## 1967	24911	21969	21956	20944	22200	24002	22951	20143	17187	15287
## 1968	26943	23735	20744	21090	21502	21275	19426	16798	14209	13357
## 1969	23460	19551	15898	16012	16054	15910	13873	11854	10138	9942
## 1970	17778	13854	12681	11328	11946	13043	12785	11937	11383	10282
## 1971	18337	16779	15504	17258	18264	19184	19453	18741	19087	18171
## 1972	37486	37303	37639	36536	35850	41581	42979	42490	37992	32454
## 1973	48622	39868	34511	37234	36675	37945	36593	31669	28682	25944
## 1974	46847	38315	32600	33349	30598	32009	37599	45999	54945	68394
## 1975	182260	184177	157547	168471	159020	160748	169631	170927	179898	176471
## 1976	248619	215342	192024	178765	182397	188423	197159	198648	195864	194125

##	1977	229415	245395	236383	226807	239984	250309	253809	254863	249551	254085
##	1978	269896	298455	290356	283308	272384	286091	290718	285424	284642	279874
##	1979	341877	357463	334400	332572	318905	311232	310000	303800	299566	286241
##	1980	334495	334265	316776	309300	308989	311232	313943	303555	290386	283822
##	1981	339700	347400	325500	315200	314900	314500	313700	318500	306000	299500
##	1982	351425	372288	358536	356004	375626	390664	404840	421856	446341	465959
##	1983	601931	632837	622819	622162	633272	635002	634020	622103	610379	599100
##	1984	674424	667059	626653	602100	600344	584506	580347	570553	565348	555279
##	1985	636841	636342	599092	580700	568574	561400	553644	541022	534700	522587
##	1986	609987	603156	578700	568400	569966	569761	573989	573735	566245	556055
##	1987	623079	619978	592892	582102	561698	550850	536522	525650	515893	492248
##	1988	517127	511023	493993	483400	481469	475070	472806	458767	441201	428578
##	1989	448572	441100	409708	393323	391918	390001	383839	377968	368060	360246
##	1990	385727	398961	390149	391108	411171	427931	441335	450824	452304	457658
##	1991	567249	580777	596890	616326	647415	676706	701677	709801	718748	720754
##	1992	779868	816124	818102	826297	838390	851831	856505			


```
## 1989 357443 374530
## 1990 480083 523798
## 1991 730105 751348
## 1992
```

#Now taking the Summary of the Dole as to check the mean median and other IQR
`summary(dole)`

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      4742   20400   56760   221200  392600   856500
```

#plotting Dole

```
plot(dole , sub="Pre-transformation")
```




#Checking the Lambda for the Dole with the help of Box Cox

```
Boxcox_lambda_dole <- BoxCox.lambda(dole)
```

#based on the Lambda we will try to plot the Boxcox transformed data

```
plot(BoxCox(dole,Boxcox_lambda_dole),sub="Post-transformation")
```

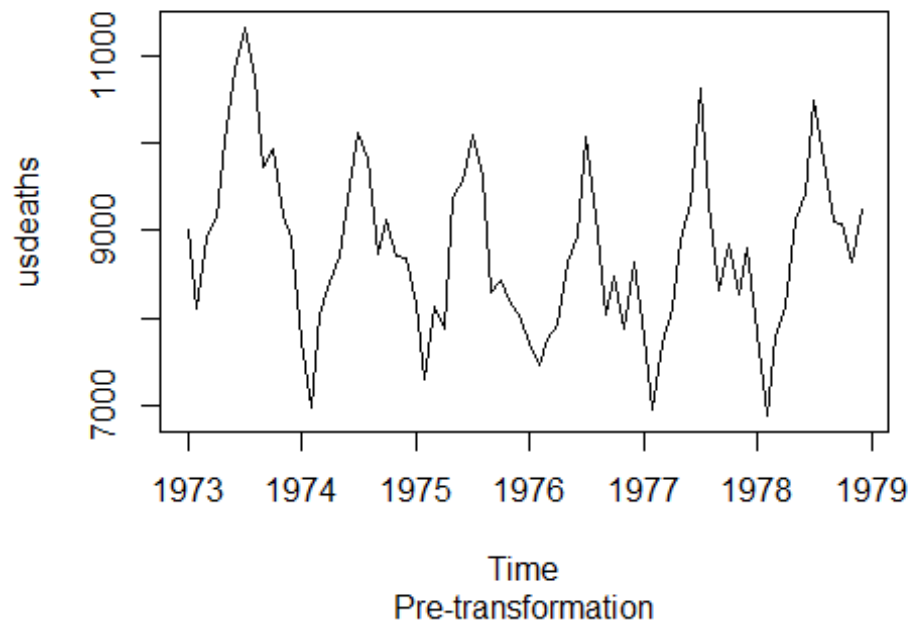


#Question 2) Monthly total of accidental deaths in the United States (January 1973-December 1978).

#Loading the Package
usdeaths

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov
## 1973	9007	8106	8928	9137	10017	10826	11317	10744	9713	9938	9161
## 1974	7750	6981	8038	8422	8714	9512	10120	9823	8743	9129	8710
## 1975	8162	7306	8124	7870	9387	9556	10093	9620	8285	8433	8160
## 1976	7717	7461	7776	7925	8634	8945	10078	9179	8037	8488	7874
## 1977	7792	6957	7726	8106	8890	9299	10625	9302	8314	8850	8265
## 1978	7836	6892	7791	8129	9115	9434	10484	9827	9110	9070	8633
##	Dec										
## 1973	8927										
## 1974	8680										
## 1975	8034										
## 1976	8647										
## 1977	8796										
## 1978	9240										

#plotting the US deaths
`plot(usdeaths , sub="Pre-transformation")`



#Taking the Summary of the US Deaths

```
summary(usdeaths)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      6892   8089   8728   8788   9323   11320
```

#Checking the Lambda for the us deaths with the help of Box Cox

```
Boxcox_lambda_usdeaths <- BoxCox.lambda(usdeaths)
```

#based on the Lambda we will try to plot the Boxcox transformed data

```
plot(BoxCox(dole,Boxcox_lambda_usdeaths),sub="Post-transformation")
```



#Question 3) Quarterly production of bricks (in millions of units) at Portland, Australia (March 1956-September 1994).


#Loading the Package
bricksq

##	Qtr1	Qtr2	Qtr3	Qtr4
## 1956	189	204	208	197
## 1957	187	214	227	223
## 1958	199	229	249	234
## 1959	208	253	267	255
## 1960	242	268	290	277
## 1961	241	253	265	236
## 1962	229	265	275	258
## 1963	231	263	308	313
## 1964	293	328	349	340
## 1965	309	349	366	340
## 1966	302	350	362	337
## 1967	326	358	359	357
## 1968	341	380	404	409
## 1969	383	417	454	428
## 1970	386	428	434	417
## 1971	385	433	453	436
## 1972	399	461	476	477
## 1973	452	461	534	516
## 1974	478	526	518	417
## 1975	340	437	459	449

```
## 1976 424 501 540 533
## 1977 457 513 522 478
## 1978 421 487 470 482
## 1979 458 526 573 563
## 1980 513 551 589 564
## 1981 519 581 581 578
## 1982 500 560 512 412
## 1983 303 409 420 413
## 1984 400 469 482 484
## 1985 447 507 533 503
## 1986 443 503 505 443
## 1987 415 485 495 458
## 1988 427 519 555 539
## 1989 511 572 570 526
## 1990 472 524 497 460
## 1991 373 436 424 430
## 1992 387 413 451 420
## 1993 394 462 476 443
## 1994 421 472 494
```

#plotting the Bricks

```
plot(bricksq ,sub="Pre-transformation")
```




#Taking the Summary of the Bricks

```
summary(bricksq)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    187.0   338.5   428.0   408.8   490.5   589.0
```

```
#Checking the Lambda for the Bricks with the help of Box Cox
Boxcox_lambda_Bricks <- BoxCox.lambda(usdeaths)
#based on the Lambda we will try to plot the Boxcox transformed data
plot(BoxCox(dole,Boxcox_lambda_Bricks) , sub="Post-transformation")
```




```
#####
#####
```

```
#Use the Dow Jones index (data set dowjones) to do the following:
```

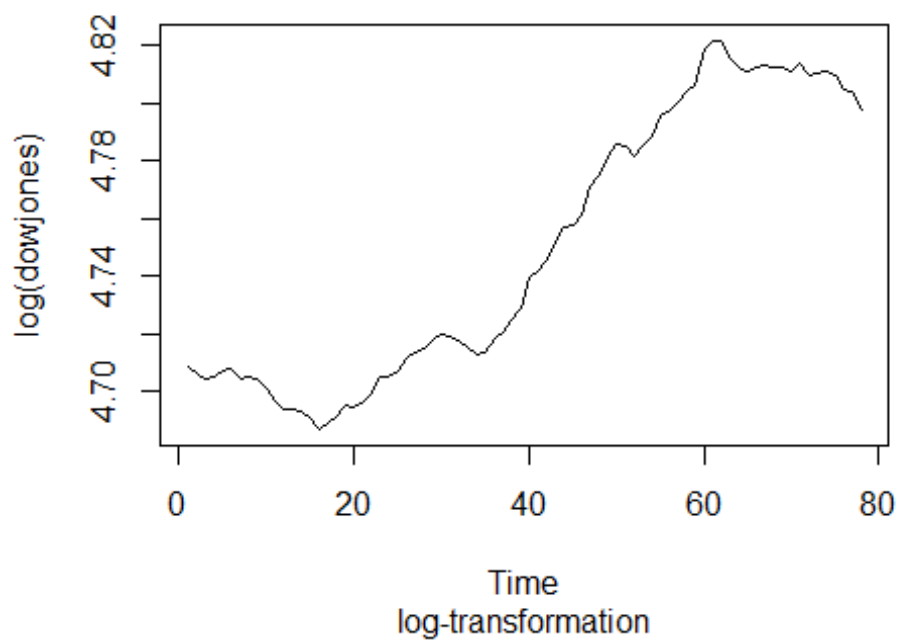
```
# Question 1 ) Produce a time plot of the series.
```

```
#plotting the Dow jones
```


```
plot(dowjones)
```

```
#ploting the log of the DOW jones  
plot(log(dowjones) , sub="log-transformation")
```



```
#ploting the log of the Square root
plot(sqrt(dowjones) , sub="Sqrt-transformation")
```

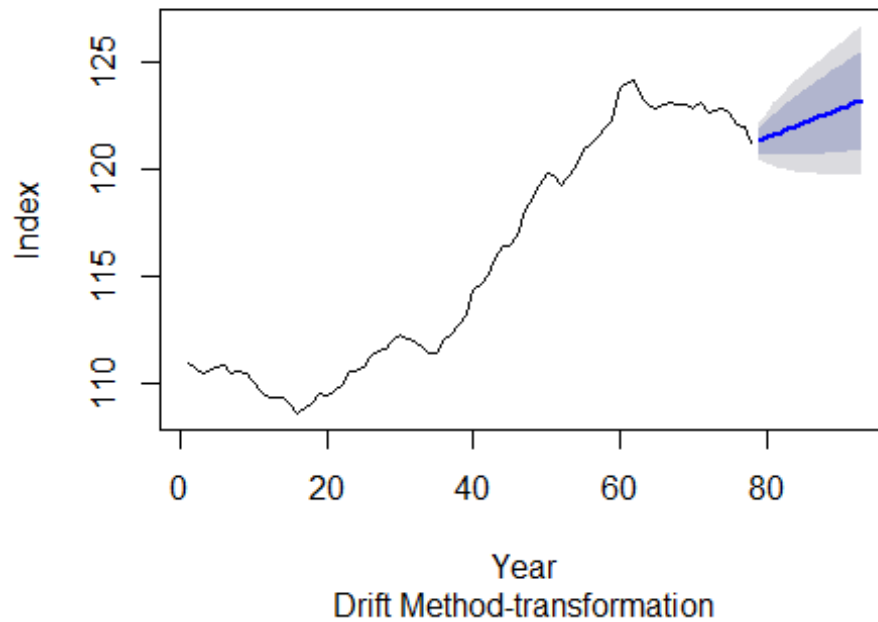


#Question 2) Produce forecasts using the drift method and plot them

```
#Generating Random Walk Forecasts using drift menthod
drift_dowjones <- rwf(dowjones , h=15, drift=TRUE)
#Generating Random Walk Forecasts Log using drift menthod
drift_dowjones_log <- rwf(log(dowjones), h=15, drift=TRUE)
#Generating Random Walk Forecasts sqrt using drift menthod
drift_dowjones_sqrt <- rwf(sqrt(dowjones), h=15, drift=TRUE)

#Plotting the Data using Drift menthod
plot(drift_dowjones , sub="Drift Method-transformation"
,ylab="Index",xlab="Year")
```

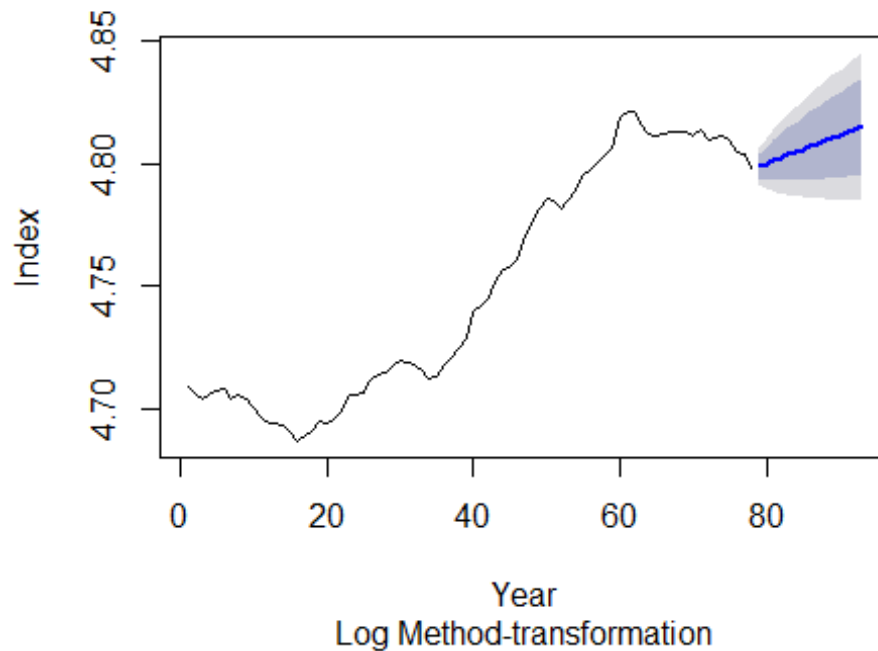
Forecasts from Random walk with drift



#Plotting the Data using Log method


```
plot(drift_dowjones_log , sub="Log Method-transformation"  
 ,ylab="Index",xlab="Year")
```

Forecasts from Random walk with drift



```
#Plotting the Data using SQRT method
plot(drift_dowjones_sqrt , sub="SQRT Method-transformation",
,ylab="Index",xlab="Year")
```

Forecasts from Random walk with drift



#Question 2) Show that the graphed forecasts are identical to extending the line drawn between the first and last observations.

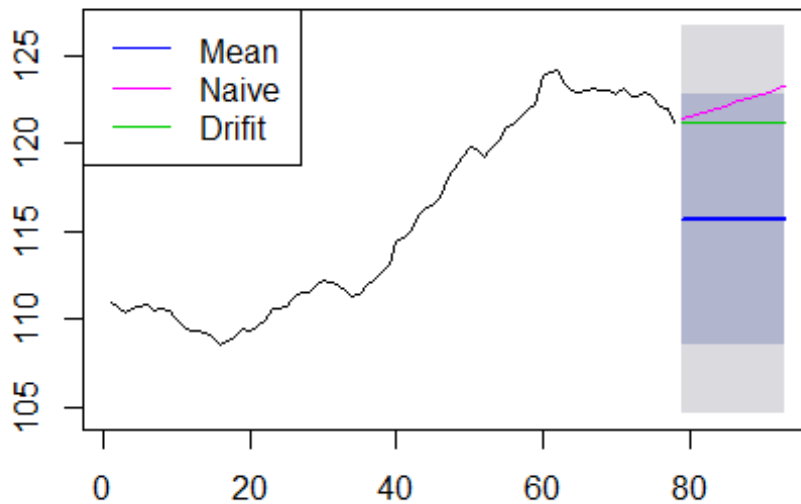
#for the First and Last observation we need to do the Mean , Naive and Drift method in the Plot

```
dowjones_win <- window(dowjones)
dowjones_mean <- meanf(dowjones_win,h=15)
#Generating Random Walk Forecasts with Drift = TRUE
dowjones_drift <- rwf(dowjones_win,h=15, drift = TRUE)
#Generating Random Walk Forecasts without Drift
dowjones_no_drift <- rwf(dowjones_win,h=15)
```

Generating the Plot for the same

```
plot(dowjones_mean, main="Dow Jones")
lines(dowjones_drift$mean,col=6)
lines(dowjones_no_drift$mean,col=3)
lines(dowjones)
legend("topleft", lty=1, col=c(4,6,3), legend=c("Mean ", "Naive", "Drift"))
```

Dow Jones



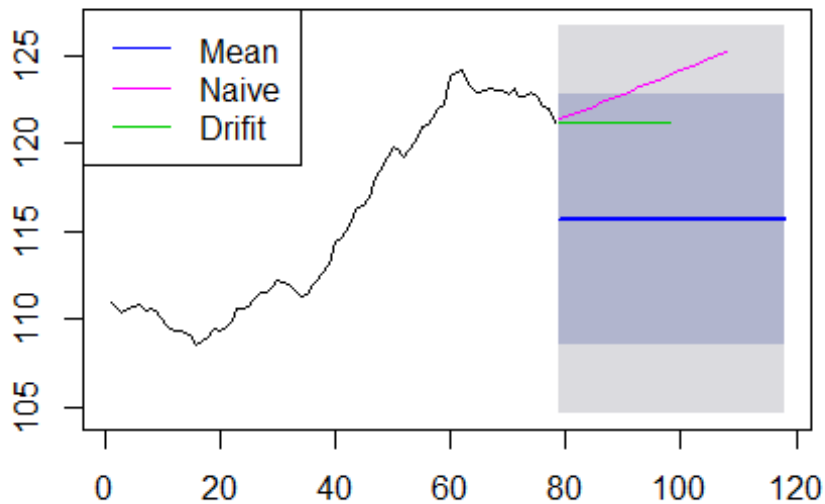
#Question 3) Try some of the other benchmark functions to forecast the same data set.

#Which do you think is best? Why?

```
dowjones_nw_win <- window(dowjones)
dowjones_nw_mean <- meanf(dowjones_nw_win,h=40)
#Generating Random Walk Forecasts with Drift = TRUE
dowjones_nw_drift <- rwf(dowjones_nw_win,h=30, drift = TRUE)
#Generating Random Walk Forecasts without Drift
dowjones_no_nw_drift <- rwf(dowjones_nw_win,h=20)

# Generating the Plot for the same with modified h
plot(dowjones_nw_mean, main="Dow Jones")
lines(dowjones_nw_drift$mean,col=6)
lines(dowjones_no_nw_drift$mean,col=3)
lines(dowjones)
legend("topleft", lty=1, col=c(4,6,3), legend=c("Mean ", "Naive", "Drift"))
```

Dow Jones



```
#####
#####
```

#Consider the daily closing IBM stock prices (data set ibmclose).

#Question 1) Produce some plots of the data in order to become familiar with it.

#Loading the Package

ibmclose

Time Series:

Start = 1

End = 369

Frequency = 1


```
## [1] 460 457 452 459 462 459 463 479 493 490 492 498 499 497 496 490 489
## [18] 478 487 491 487 482 479 478 479 477 479 475 479 476 476 478 479 477
## [35] 476 475 475 473 474 474 474 465 466 467 471 471 467 473 481 488 490
## [52] 489 489 485 491 492 494 499 498 500 497 494 495 500 504 513 511 514
## [69] 510 509 515 519 523 519 523 531 547 551 547 541 545 549 545 549 547
## [86] 543 540 539 532 517 527 540 542 538 541 541 547 553 559 557 557 560
## [103] 571 571 569 575 580 584 585 590 599 603 599 596 585 587 585 581 583
## [120] 592 592 596 596 595 598 598 595 595 592 588 582 576 578 589 585 580
## [137] 579 584 581 581 577 577 578 580 586 583 581 576 571 575 575 573 577
## [154] 582 584 579 572 577 571 560 549 556 557 563 564 567 561 559 553 553
```

```
## [171] 553 547 550 544 541 532 525 542 555 558 551 551 552 553 557 557 548
## [188] 547 545 545 539 539 535 537 535 536 537 543 548 546 547 548 549 553
## [205] 553 552 551 550 553 554 551 551 545 547 547 537 539 538 533 525 513
## [222] 510 521 521 521 523 516 511 518 517 520 519 519 519 518 513 499 485
## [239] 454 462 473 482 486 475 459 451 453 446 455 452 457 449 450 435 415
## [256] 398 399 361 383 393 385 360 364 365 370 374 359 335 323 306 333 330
## [273] 336 328 316 320 332 320 333 344 339 350 351 350 345 350 359 375 379
## [290] 376 382 370 365 367 372 373 363 371 369 376 387 387 376 385 385 380
## [307] 373 382 377 376 379 386 387 386 389 394 393 409 411 409 408 393 391
## [324] 388 396 387 383 388 382 384 382 383 383 388 395 392 386 383 377 364
## [341] 369 355 350 353 340 350 349 358 360 360 366 359 356 355 367 357 361
## [358] 355 348 343 330 340 339 331 345 352 346 352 357
```


#Now taking the Summary of the ibm as to check the mean median and other IQR
summary(ibmclose)

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##    306.0   387.0   494.0   478.5   549.0   603.0
```


#Producing some plots
Generating the Plot Log
plot(log(ibmclose))




Generating the Plot SQRT
plot(sqrt(ibmclose))



```
# Generating the Plot Simple  
plot(ibmclose)
```




```
## Generating the Plot QQ NORM PLOT  
qqnorm(ibmclose)
```



#Question 2) Split the data into a training set of 300 observations and a test set of 69 observations.

```
ibmclose_train_data <- window(ibmclose ,Start= 1 ,end=300)  
ibmclose_test_data <- window(ibmclose ,start=301, end=369)
```

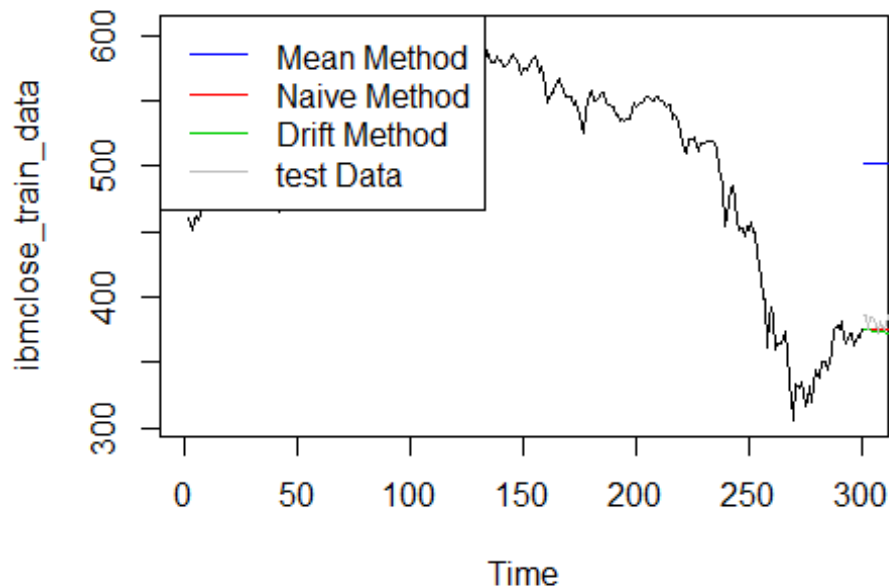
#Question 3) Try various benchmark methods to forecast the training set and compare the results on the test set. Which method did best?

```
ibmclose_avg <- meanf(ibmclose_train_data,h=64)$mean  
ibmclose_naive <- naive(ibmclose_train_data ,h=64)$mean  
ibmclose_drift <- rwf(ibmclose_train_data ,drift=TRUE,h=64)$mean
```

#plotting the Data

```
plot(ibmclose_train_data)  
lines(ibmclose_naive,col=2)  
lines(ibmclose_avg,col=4)  
lines(ibmclose_drift,col=3)  
lines(ibmclose_test_data,col=8)
```

```
legend("topleft",lty=1,col=c(4,2,3,8), legend=c("Mean Method","Naive
Method","Drift Method","test Data"))
```



```
#####
#####
```

#Consider the sales of new one-family houses in the USA, Jan 1973 - Nov 1995 (data set hsales).

#Question 1) Produce some plots of the data in order to become familiar with it.

#Loading the Package
hsales

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1973  55  60  68  63  65  61  54  52  46  42  37  30
## 1974  37  44  55  53  58  50  48  45  41  34  30  24
## 1975  29  34  44  54  57  51  51  53  46  46  46  39
## 1976  41  53  55  62  55  56  57  59  58  55  49  47
## 1977  57  68  84  81  78  74  64  74  71  63  55  51
## 1978  57  63  75  85  80  77  68  72  68  70  53  50
## 1979  53  58  73  72  68  63  64  68  60  54  41  35
```

```
## 1980 43 44 44 36 44 50 55 61 50 46 39 33
## 1981 37 40 49 44 45 38 36 34 28 29 27 29
## 1982 28 29 36 32 36 34 31 36 39 40 39 33
## 1983 44 46 57 59 64 59 51 50 48 51 45 48
## 1984 52 58 63 61 59 58 52 48 53 55 42 38
## 1985 48 55 67 60 65 65 63 61 54 52 51 47
## 1986 55 59 89 84 75 66 57 52 60 54 48 49
## 1987 53 59 73 72 62 58 55 56 52 52 43 37
## 1988 43 55 68 68 64 65 57 59 54 57 43 42
## 1989 52 51 58 60 61 58 62 61 49 51 47 40
## 1990 45 50 58 52 50 50 46 46 38 37 34 29
## 1991 30 40 46 46 47 47 43 46 37 41 39 36
## 1992 48 55 56 53 52 53 52 56 51 48 42 42
## 1993 44 50 60 66 58 59 55 57 57 56 53 51
## 1994 45 58 74 65 65 55 52 59 54 57 45 40
## 1995 47 47 60 58 63 64 64 63 55 54 44
```

#Now taking the Summary of the ibm as to check the mean median and other IQR
`summary(hsales)`

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##  24.00   44.00   53.00   52.29   59.00   89.00
```


#Producing some plots

Generating the Plot Log


`plot(log(hsales))`

Generating the Plot Log


`plot(log(hsales))`



```
# Generating the Plot SQRT  
plot(sqrt(hsales))
```




```
# Generating the Plot Simple  
plot(hsales)
```



```
## Generating the Plot QQ NORM PLOT  
qqnorm(hsales)
```

Normal Q-Q Plot



#Question 2) Split the hsales data set into a training set and a test set, where the test set is the last two years of data.

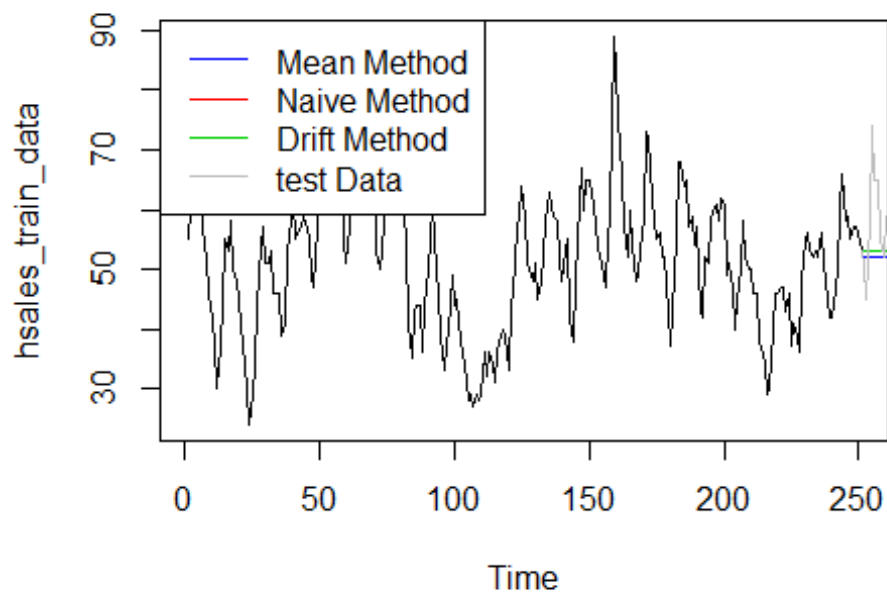
```
hsales_ts <- ts(hsales,start=1,end=275)
hsales_train_data <- window(hsales_ts,start=1,end=251)
hsales_test_data <- window(hsales_ts,start=251)
```

#Question 3) Try various benchmark methods to forecast the training set and compare the results on the test set. Which method did best?

```
hsales_avg <- meanf(hsales_train_data,h=34)$mean
hsales_naive <- naive(hsales_train_data ,h=34)$mean
hsales_drift <- rwf(hsales_train_data ,drift=TRUE,h=34)$mean
```

#plotting the Data

```
plot(hsales_train_data)
lines(hsales_naive,col=2)
lines(hsales_avg,col=4)
lines(hsales_drift,col=3)
lines(hsales_test_data,col=8)
legend("topleft",lty=1,col=c(4,2,3,8), legend=c("Mean Method", "Naive
Method", "Drift Method", "test Data"))
```



```
#####
#####
#
##
#
##
#
##
#####
#####

#Day      1    2    3    4    5    6    7    8    9    10   11   12
#Mwh      16.3  16.8  15.5  18.2  15.2  17.5  19.8  19.0  17.5
16.0      19.6  18.0
#temp     29.3  21.7  23.7  10.4  29.7  11.9  9.0  23.4  17.8
30.0      8.6 11.8

#Question 1 ) Plot the data and find the regression model for Mwh with
temperature as an explanatory variable. Why is there a negative relationship?

#Loading the Package
econsumption


##      Mwh temp
## 1  16.3 29.3
```

```
## 2 16.8 21.7
## 3 15.5 23.7
## 4 18.2 10.4
## 5 15.2 29.7
## 6 17.5 11.9
## 7 19.8 9.0
## 8 19.0 23.4
## 9 17.5 17.8
## 10 16.0 30.0
## 11 19.6 8.6
## 12 18.0 11.8
```

#Now taking the Summary of the ibm as to check the mean median and other IQR
`summary(econsumption)`


```
##      Mwh      temp
##  Min.   :15.20  Min.   : 8.60
##  1st Qu.:16.23  1st Qu.:11.45
##  Median :17.50  Median :19.75
##  Mean   :17.45  Mean   :18.94
##  3rd Qu.:18.40  3rd Qu.:25.10
##  Max.   :19.80  Max.   :30.00
```

#Producing some plots
Generating the Plot Log
`plot(log(econsumption))`





#question 2) a. Plot the data and find the regression model for Mwh with temperature as an explanatory variable. Why is there a negative relationship?



```
plot(Mwh ~ temp, data = econsumption)
fit_eco = lm(formula = Mwh ~ temp, data = econsumption)
abline(fit_eco, col=5)
```



#Question 3) Produce a residual plot. Is the model adequate? Are there any outliers or influential observations?

```
plot(fit_eco)
```





#Question 4) Use the model to predict the electricity consumption that you would expect for a day with maximum temperature 10 and a day with maximum temperature 35. Do you believe these predictions?

```
coeffs = coefficients(fit_eco)
```


```
pred_temp = c(10, 35)
p_temp = coeffs[1] + coeffs[2]*pred_temp
p_temp
```

```
## [1] 18.74795 15.11902
```

#Question 5 Give prediction intervals for your forecasts. The following R code will get you started:


```
forecast <- forecast(fit_eco , newdata=data.frame(temp=10))
plot(forecast)
```

Forecasts from Linear regression model



```
forecast2 <- forecast(fit_eco, newdata=data.frame(temp=35))
plot(forecast2)
```

Forecasts from Linear regression model



```
temp10 = data.frame(temp=10)
temp35 = data.frame(temp=35)
predict(fit_eco, temp10, interval="predict")
```

```
##      fit      lwr      upr
## 1 18.74795 16.34824 21.14766
```

```
predict(fit_eco, temp35, interval="predict")
```

```
##      fit      lwr      upr
## 1 15.11902 12.49768 17.74035
```

```
#####
#####
```

#The following table gives the winning times (in seconds) for the men's 400 meters final in each Olympic Games from 1896 to 2012 (data set `olympic`).

#question 1) Update the data set `olympic` to include the winning times from the last few Olympics.

```
olympic
```

```
##   Year  time
## 1 1896 54.20
## 2 1900 49.40
## 3 1904 49.20
```


```
## 4 1908 50.00
## 5 1912 48.20
## 6 1920 49.60
## 7 1924 47.60
## 8 1928 47.80
## 9 1932 46.20
## 10 1936 46.50
## 11 1948 46.20
## 12 1952 45.90
## 13 1956 46.70
## 14 1960 44.90
## 15 1964 45.10
## 16 1968 43.80
## 17 1972 44.66
## 18 1976 44.26
## 19 1980 44.60
## 20 1984 44.27
## 21 1988 43.87
## 22 1992 43.50
## 23 1996 43.49
```

#question 2) Plot the winning time against the year. Describe the main features of the scatterplot.

```
plot(time ~ Year, data = olympic)
```

#Question 3) Fit a regression line to the data. Obviously the winning times have been decreasing, but at what average rate per year?



```
oly_fit = lm(formula = time ~ Year, data = olympic)
plot(time ~ Year, data = olympic)
abline(oly_fit, col=5)
```

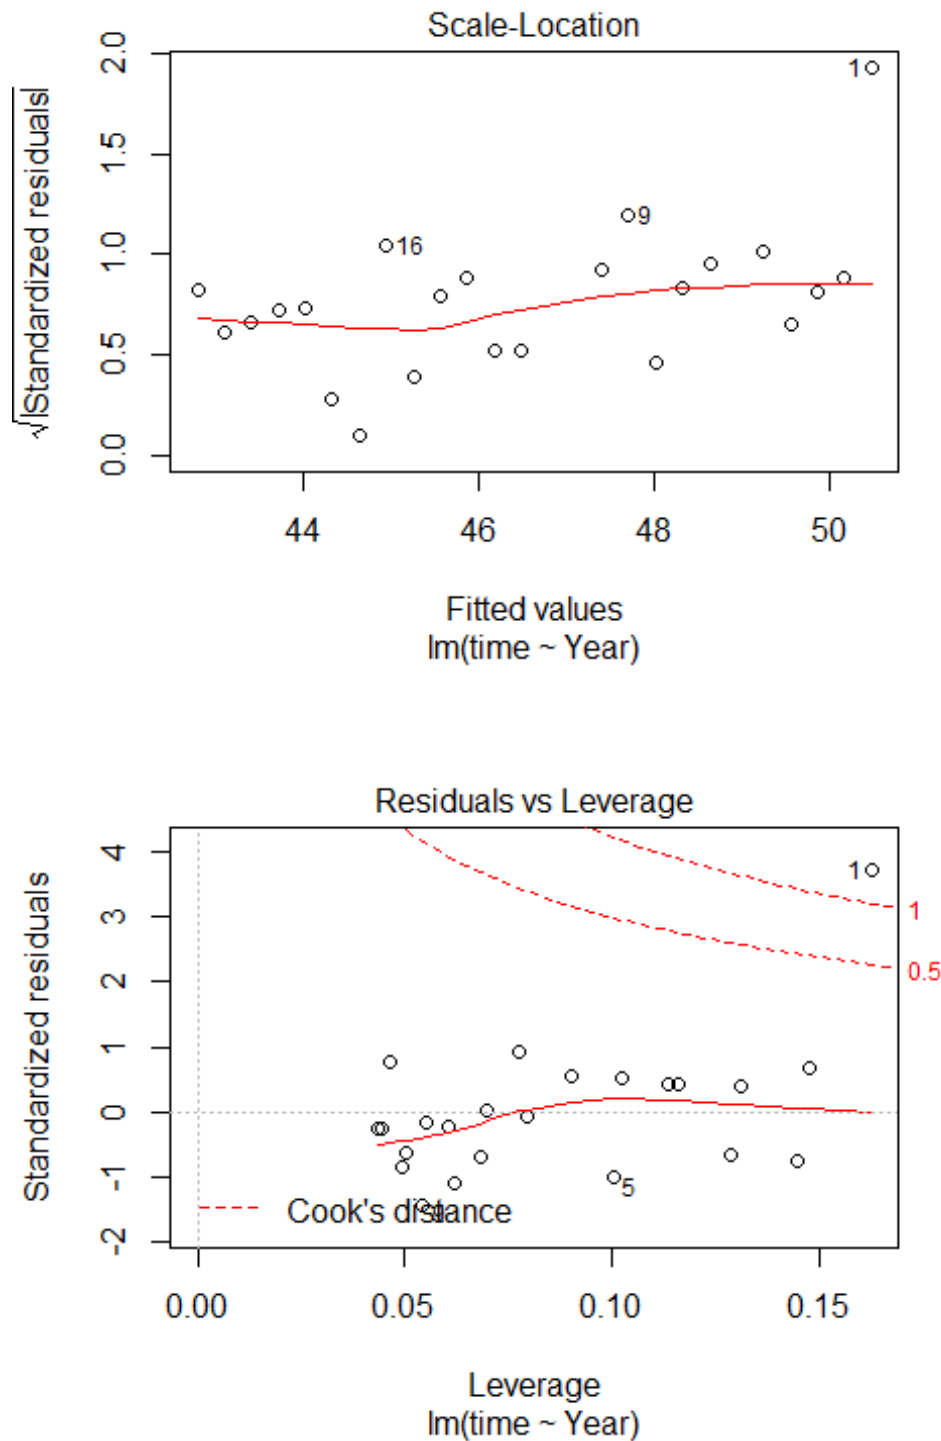


```
olympic_ts <- ts(olympic,start=1,end=28)
```

#Question 4) Plot the residuals against the year. What does this indicate about the suitability of the fitted line?

```
plot(oly_fit)
```





#Question 5) Predict the winning time for the men's 400 meters final in the 2000, 2004, 2008 and 2012 Olympics. Give a prediction interval for each of your forecasts. What assumptions have you made in these calculations?

```


coeff = coefficients(oly_fit)
pred_time = c(2000, 2004, 2008, 2012)
p_time = coeff[1] + coeff[2]*pred_time
p_time

## [1] 42.49977 42.19261 41.88545 41.57829

fcast3 <- forecast(oly_fit, newdata=data.frame(Year=2000))
plot(fcast3, xlab="Year", ylab="time")

```

Forecasts from Linear regression model




```

fcast4 <- forecast(oly_fit, newdata=data.frame(Year=2004))
plot(fcast4, xlab="Year", ylab="time")


```

Forecasts from Linear regression model




```
fcast5 <- forecast(oly_fit, newdata=data.frame(Year=2008))  
plot(fcast5, xlab="Year", ylab="time")
```

Forecasts from Linear regression model



```
fcast6 <- forecast(oly_fit, newdata=data.frame(Year=2012))
plot(fcast6, xlab="Year", ylab="time")
```

Forecasts from Linear regression model



```
time2000 = data.frame(Year=2000)
time2004 = data.frame(Year=2004)
time2008 = data.frame(Year=2008)
time2012 = data.frame(Year=2012)

predict(oly_fit, time2000, interval="predict")

##          fit          lwr          upr
## 1 42.49977 40.05401 44.94554

predict(oly_fit, time2004, interval="predict")

##          fit          lwr          upr
## 1 42.19261 39.72657 44.65866

predict(oly_fit, time2008, interval="predict")

##          fit          lwr          upr
## 1 41.88545 39.39782 44.37308

predict(oly_fit, time2012, interval="predict")

##          fit          lwr          upr
## 1 41.57829 39.0678 44.08879
```

#Question 6) Find out the actual winning times for these Olympics (see www.databaseolympics.com). How good were your forecasts and prediction intervals?

```
coeffs1 = coefficients(oly_fit)
pred_time = c(2000, 2004, 2008, 2012)
p_time = coeffs1[1] + coeffs1[2]*pred_time
p_time
```

```
## [1] 42.49977 42.19261 41.88545 41.57829
```

```
#####
#####
```

#An elasticity coefficient is the ratio of the percentage change in the forecast variable (yy) to the percentage change in the predictor variable (xx). Mathematically, the elasticity is defined as $(dy/dx) \times (x/y)$. Consider the Log-Log model,

$\log y = \beta_0 + \beta_1 \log x + \epsilon$.

#Express yy as a function of xx and show that the coefficient β_1 is the elasticity coefficient.

*# $\log y = \beta_0 + \beta_1 \log x + \epsilon$
#Taking Differential*

$dy/y = dx/x \cdot \beta_1$ ($dy/y / dx/x$) = β_1

change in y divided by y over the the change in x divided by x

$100(dy/y) = 100(dx/x)\beta_1$ % change y = % change x β_1

```
#####
#####
```


###The data below (data set fancy) concern the monthly sales figures of a shop which opened in January 1987 and sells gifts, souvenirs, and novelties. The shop is situated on the wharf at a beach resort town in Queensland, Australia. The sales volume varies with the seasonal population of tourists. There is a large influx of visitors to the town at Christmas and for the local surfing festival, held every March since 1988. Over time, the shop has expanded its premises, range of products, and staff.

#Question 1)) Produce a time plot of the data and describe the patterns in the graph.

#Identify any unusual or unexpected fluctuations in the time series.

```
library(fma)
library(fpp)

## Warning: package 'fpp' was built under R version 3.3.3
## Loading required package: expsmooth
## Warning: package 'expsmooth' was built under R version 3.3.3
## Loading required package: lmtest
## Warning: package 'lmtest' was built under R version 3.3.3
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
## Loading required package: tseries
## Warning: package 'tseries' was built under R version 3.3.3
plot(fancy)
```



#Question 2) Explain why it is necessary to take Logarithms of these data before fitting a model.

Answer 2) It is because of the increasing seasonal fluctuations
`log_fancy <- log(fancy)`

#question 3) c) Use R to fit a regression model to the Logarithms of these sales data with a linear trend, seasonal dummies and a "surfing festival" dummy variable.

```
log_fancy <- log(fancy)
dummy_fest = rep(0, length(fancy))
dummy_fest[seq_along(dummy_fest)%%12 == 3] <- 1
dummy_fest[3] <- 0
dummy_fest <- ts(dummy_fest, freq = 12, start=c(1987,1))
my_data <- data.frame(
  log_fancy,
  dummy_fest
)


fit <- tslm(log_fancy ~ trend + season + dummy_fest, data=my_data)

future_data <- data.frame(
  dummy_fest = rep(0, 12)
)
future_data[3,] <- 1
forecast(fit, newdata=future_data)
```


##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	Jan 1994	9.491352	9.238522	9.744183	9.101594	9.881111
##	Feb 1994	9.764789	9.511959	10.017620	9.375031	10.15455
##	Mar 1994	10.302990	10.048860	10.557120	9.911228	10.69475
##	Apr 1994	9.941465	9.688635	10.194296	9.551707	10.33122
##	May 1994	9.988919	9.736088	10.241749	9.599161	10.37868
##	Jun 1994	10.050280	9.797449	10.303110	9.660522	10.44004
##	Jul 1994	10.233926	9.981095	10.486756	9.844168	10.62368
##	Aug 1994	10.233456	9.980625	10.486286	9.843698	10.62321
##	Sep 1994	10.336841	10.084010	10.589671	9.947083	10.72660
##	Oct 1994	10.436923	10.184092	10.689753	10.047165	10.82668
##	Nov 1994	10.918299	10.665468	11.171129	10.528541	11.30806
##	Dec 1994	11.695812	11.442981	11.948642	11.306054	12.08557

#question 4 d) Plot the residuals against time and against the fitted values. Do these plots reveal any problems with the model?

```
plot(residuals(fit), type='p')
```




```
plot(as.numeric(fitted(fit)), residuals(fit), type='p')
```



#Question 5) e) Do boxplots of the residuals for each month. Does this reveal any problems with the model?

```
boxplot(resid(fit) ~ cycle(resid(fit)))
```



#Question 5) ; f) What do the values of the coefficients tell you about each variable?

*# The value of the coefficients show how much the model thinks each month contributes
to the conditional mean of the model.*

#Question 5 g) What does the Durbin-Watson statistic tell you about your model?

```
dwtest(fit)
```

```
##  
## Durbin-Watson test  
##  
## data: fit  
## DW = 0.88889, p-value = 9.78e-08  
## alternative hypothesis: true autocorrelation is greater than 0
```

Question 6 h) Regardless of your answers to the above questions, use your regression model to predict the monthly sales for 1994, 1995, and 1996.

Produce prediction intervals for each of your forecasts.

```
future_data <- data.frame(  
  dummy_fest = rep(0, 36)  
)  
preds <- forecast(fit, newdata=future_data)
```

Question 5 i) Transform your predictions and intervals to obtain predictions and intervals for the raw data.

```
df_pred <- as.data.frame(preds)  
df_pred <- exp(df_pred)
```

Question 5 j) How could you improve these predictions by modifying the model?

#ANSwers 5 J) We Could use consider using a dynamic-regression model

```
#####  
#####
```


Question 5b) (1) The data below (data set texasgas) shows the demand for natural gas and the price of natural gas for 20 towns in Texas in 1969.

a) Do a scatterplot of consumption against price. The data are clearly not linear. Three possible nonlinear models for the data are given below

```
library(fma)  
library(fpp)  
library(segmented)
```

```
## Warning: package 'segmented' was built under R version 3.3.2
```

```
df_gas <- (texasgas)  
plot(df_gas$price, df_gas$consumption)
```



b) Can you explain why the slope of the fitted line should change with price?

Ans The data is not linear so the slope needs to change in order to get the Data from our model

c) Fit the three models and find the coefficients, and residual variance in each case.

First Model

```
fit_gas <- lm(consumption ~ exp(price), df_gas)
```

```
fit_gas
```

```
##
```

```
## Call:
```

```
## lm(formula = consumption ~ exp(price), data = df_gas)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)    exp(price)
```

```
##  7.086e+01   -1.642e-43
```

Residual Variance

```
(summary(fit_gas)$sigma)**2
```

```
## [1] 1101.359

# Second model - piecewise linear regression

lin.mod <- lm(consumption ~ price, df_gas)

lin.mod

##
## Call:
## lm(formula = consumption ~ price, data = df_gas)
##
## Coefficients:
## (Intercept)      price
##      138.561      -1.104

segmented.mod <- segmented(lin.mod, seg.Z = ~price, psi=60)

slope(segmented.mod)

## $price
##           Est. St.Err. t value CI(95%).l CI(95%).u
## slope1 -3.1470  0.5102  -6.169   -4.2290   -2.0660
## slope2 -0.3075  0.2220  -1.385   -0.7782    0.1632

# Residual variance
(summary(segmented.mod)$sigma)**2

## [1] 167.8511

# Third model - polynomial regression

poly_fit <- lm(consumption ~ poly(price, 2), df_gas)

# Residual variance
(summary(poly_fit)$sigma)**2

## [1] 206.5276


#d) For each model, find the value of R2 and AIC, and produce a residual
plot. Comment on the adequacy of the three models.

# First model - basic linear regression

# Adjusted R-squared: -0.004
# AIC: 200.736

resid <- residuals(fit_gas)
plot(fit_gas$fitted.values, resid, ylab='residuals', xlab='fitted values',
```

```
main='linear regression')
abline(0,0)
```




```
# Second model - piecewise linear regression
```

```
# Adjusted R-squared: 0.847
```

```
# AIC: 164.756
```

```
resid <- residuals(segmented.mod)
plot(segmented.mod$fitted.values, resid, ylab='residuals', xlab='fitted
values',
      main='piecewise linear regression')
abline(0,0)
```

piecewise linear regression



```
# Third model - polynomial regression.
```

```
# Adjusted R-squared: 0.812
```


```
# AIC: 168.116
```

```
resid <- residuals(poly_fit)
```

```
plot(poly_fit$fitted.values, resid, ylab='residuals', xlab='fitted values',  
      main='polynomial linear regression')
```

```
abline(0,0)
```

polynomial linear regression



e) For prices 40, 60, 80, 100, and 120 cents per 1,000 cubic feet, compute the forecasted per capita demand using the best model of the three above.

```
new.data <- data.frame(price=c(40, 60, 80, 100, 120))
predict(segmented.mod, new.data)
```

```
##          1          2          3          4          5
## 104.53618  53.34514  47.19593  41.04673  34.89752
```

F) Compute 95% prediction intervals. Make a graph of these prediction intervals and discuss their interpretation.

```
newx <- seq(min(new.data), max(new.data), length.out=5)
intervals <- predict(segmented.mod, new.data, interval="predict")
```

```
plot(consumption ~ price, data = df_gas, type = 'n')
```

```
polygon(c(rev(newx), newx), c(rev(intervals[,3]), intervals[,2]), col =
'grey80', border = NA)
```

